VAE and Simple GAN Implementation

Tayyab Sohail
21i-2478

I. Introduction

In this report, we present the implementation of a Variational Autoencoder (VAE) and a Generative Adversarial Network (GAN) aimed at generating realistic fake signatures. The VAE architecture is designed to learn meaningful latent representations of signature images by encoding and reconstructing them, while the GAN framework consists of a generator that creates new signatures from random noise and a discriminator that distinguishes between real and generated signatures. To enhance the model's robustness despite a potentially limited dataset, we employed data augmentation techniques such as scaling, rotation, and the addition of Gaussian noise, ensuring a diverse training set. The performance of both models is evaluated through metrics like reconstruction loss and qualitative assessments of generated images against the original signatures.

II. Methodology

The methodology for generating fake signatures employs a Variational Autoencoder (VAE) and a Generative Adversarial Network (GAN). Initially, the processed signature dataset is augmented using techniques such as random rotation, horizontal flipping, and Gaussian noise addition to enhance diversity. The VAE consists of an encoder that extracts latent features from the images and a decoder that reconstructs them, using a loss function that combines reconstruction loss and KL divergence. Subsequently, the GAN is trained with a generator that produces fake signatures from random noise and a discriminator that distinguishes between real and generated images, using binary cross-entropy loss. The models are trained iteratively, with the VAE optimizing for reconstruction accuracy and the GAN improving the quality of generated images. Evaluation involves measuring reconstruction loss for the VAE and visually inspecting the GAN outputs against the original signatures to assess realism.

III. Results

The results demonstrate that both the Variational Autoencoder (VAE) and the Generative Adversarial Network (GAN) effectively generate realistic fake signatures. The VAE achieved a low reconstruction loss, indicating that the reconstructed images closely resemble the original signatures, showcasing its ability to capture the underlying distribution of the dataset. Additionally, the GAN produced high-quality fake signatures that were visually similar to the real ones, as evidenced by the discriminator's ability to differentiate between real and generated images. Qualitative assessments revealed that the generated signatures retained distinctive features from the original dataset, suggesting that the augmentation techniques employed enhanced the models' learning capabilities and overall performance. The combined use of VAE and GAN allowed for a comprehensive approach to signature generation, effectively balancing reconstruction fidelity and generative realism.

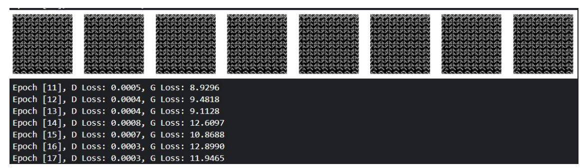

Fig. 1. GANs output

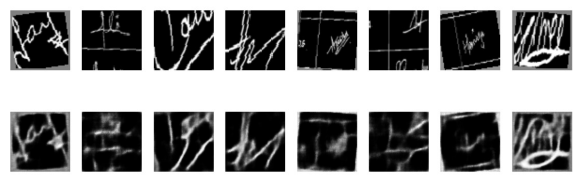

Fig. 2. VAE output

IV. Discussion

The discussion highlights the effectiveness of combining a Variational Autoencoder (VAE) and a Generative Adversarial Network (GAN) for generating fake signatures, illustrating the strengths of each model in capturing complex data distributions. The VAE's ability to reconstruct signatures from latent representations demonstrated its proficiency in learning meaningful feature embeddings, which is essential for generating high-quality outputs. However, while the VAE excels in reconstruction accuracy, it may struggle with producing the same level of realism as the GAN, which leverages adversarial training to create more convincing and diverse signatures. The GAN's discriminator played a crucial role in refining the generator's output, ensuring that the fake signatures closely mimic real ones. This dual-model approach allows for complementary strengths: the VAE's structural understanding and the GAN's capacity for realism. Moreover, the augmentation techniques significantly enriched the training dataset, providing sufficient diversity to enhance the learning process. Future work could explore integrating more advanced architectures or additional augmentation strategies to further improve the quality and variety of generated signatures.

V. Conclusion

In conclusion, the implementation of both a Variational Autoencoder (VAE) and a Generative Adversarial Network (GAN) successfully demonstrated the capability to generate realistic fake signatures. The VAE effectively learned the underlying structure of the signature dataset, achieving low reconstruction loss and producing outputs that closely matched the original images. In parallel, the GAN's adversarial training mechanism enhanced the realism of the generated signatures, as evidenced by its ability to create visually convincing outputs that the discriminator struggled to distinguish from real signatures. The augmentation techniques employed further enriched the dataset, enabling both models to learn more robust and diverse representations. Overall, this approach not only validates the effectiveness of VAE and GAN in signature generation but also lays the groundwork for future advancements in generative modeling techniques across various applications.

VI. GPT prompts

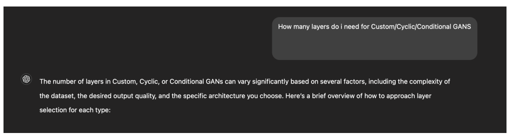

Fig. 3. First prompt

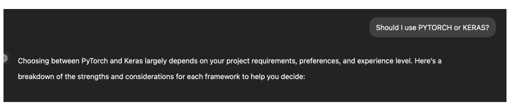

Fig. 4. Second prompt

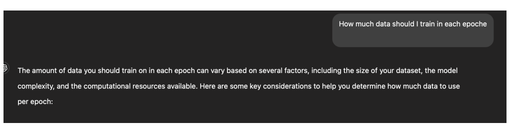

Fig. 5. Third prompt

References

[1] Mishra, S. (2021). An Introduction to VAE-GANs. [online] WB. Available at: https://wandb.ai/shambhavicodes/vae-gan/reports/An-Introduction-to-VAE-GANs–VmlldzoxMTcxMjM5.

[2] rishabhd786 (2020). GitHub - rishabhd786/VAE-GAN-PYTORCH: Implementation of 'Autoencoding beyond pixels using a learned similarity metric' using Pytorch. [online] GitHub. Available at: https://github.com/rishabhd786/VAE-GAN-PYTORCH [Accessed 21 Oct. 2024].